

## Q1e. Normalized Training and Test Error vs. Lambda

```
[3]: rng=np.random.default_rng()
train_n=100
test_n=1000
d=100
mu=np.zeros(d)

[4]: #Generate random dxd covariance matrix
A=np.random.rand(d,d)
while np.linalg.matrix_rank(A)!=d:
    A=np.random.rand(d,d)
Cov=A.T@A

sigma_noise=rng.uniform(0.3,0.7)

[5]: #Generate training and test data
X_train=rng.multivariate_normal(mu,Cov,size=train_n)
a_true=rng.normal(0,1,size=(d,1))
y_train=X_train.dot(a_true)+np.random.normal(0,sigma_noise,size=(train_n,1))
X_test=rng.multivariate_normal(mu,Cov,size=test_n)
y_test=X_test.dot(a_true)+np.random.normal(0,sigma_noise,size=(test_n,1))

[13]: lambda_values = [0.0001, 0.001, 0.01, 0.1, 1, 10, 100]
avg_train_errors = []
avg_test_errors = []
num_trials = 30

•[14]: for lambda_val in lambda_values:
    train_error_sum = 0
    test_error_sum = 0

    for _ in range(num_trials):
        # Generate new training and test data with fixed parameters
        # Calculate the ridge regression solution
        XTX = np.dot(X_train.T, X_train)
        w = np.linalg.solve(XTX + lambda_val * np.identity(d), np.dot(X_train.T, y_train))

        # Normalized training error
        train_error = np.linalg.norm(np.dot(X_train, w) - y_train) / np.linalg.norm(y_train)


        # Normalized test error
        test_error = np.linalg.norm(np.dot(X_test, w) - y_test) / np.linalg.norm(y_test)

        train_error_sum += train_error
        test_error_sum += test_error

    # Average errors across trials for this lambda
    avg_train_error = train_error_sum / num_trials
    avg_test_error = test_error_sum / num_trials
    avg_train_errors.append(avg_train_error)
    avg_test_errors.append(avg_test_error)
```

```
[15]: plt.grid(True)
plt.semilogx(lambda_values, avg_train_errors, label='Normalized Training Error')
plt.semilogx(lambda_values, avg_test_errors, label='Normalized Test Error')
plt.xlabel('Lambda (log scale)')
plt.ylabel('Normalized Error')
plt.title('Normalized Training and Test Error vs. Lambda')
plt.legend()
```

[15]: <matplotlib.legend.Legend at 0x22dcda73e50>



## Q2. Loading the Data

```
[77]: import numpy as np
import torch, torchvision
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
```

```
[133]: train_set = torchvision.datasets.FashionMNIST("./data", download=True)
test_set = torchvision.datasets.FashionMNIST("./data", download=True, train=False)
X_train = train_set.data.numpy()
labels_train = train_set.targets.numpy()
X_test = test_set.data.numpy()
labels_test = test_set.targets.numpy()
```

```
[134]: X_train = X_train.reshape((X_train.shape[0], X_train.shape[1]*X_train.shape[2]))
X_test = X_test.reshape((X_test.shape[0], X_test.shape[1]*X_test.shape[2]))
```

```
[135]: X_train = X_train/255.0
X_test = X_test/255.0
```

## Q2b. Function to train the classifier

```
[155]: def train(X, Y, lambda_):  
        d = X.shape[1]  
        k = Y.shape[1] # Number of classes  
        W = np.linalg.solve(X.T @ X + lambda_ * np.identity(d), X.T @ Y)  
        return W
```

### Function to predict labels

```
[137]: def predict(W, X):  
        return np.argmax(X @ W, axis=1)
```

```
[156]: lambda_ = 1e-4  
        num_classes = 10  
        y = np.eye(num_classes)[labels_train]  
  
        # Train  
        W_hat = train(X_train, y, lambda_)
```

```
[157]: predicted_labels_test = predict(W_hat, X_test)
```

```
[158]: def evaluate_prediction(predictions, true_labels):  
        num_errors = np.sum(predictions != true_labels)  
        error_rate = num_errors / len(true_labels)  
        return error_rate
```

```
[160]: predicted_labels_train = predict(W_hat, X_train)  
  
        # Calculate training error  
        training_error = evaluate_prediction(predicted_labels_train, labels_train)  
        print("Training Error:", training_error)
```

Training Error: 0.17538333333333334

```
[159]: predicted_labels_train = predict(W_hat, X_train)  
  
        # Calculate test error  
        training_error = evaluate_prediction(predicted_labels_test, labels_test)  
        print("Testing Error:", training_error)
```

Testing Error: 0.1913

## Q2c. Partitioning and Plotting $\hat{W}^p$ vs $p$

```
[148]: # Parameters
num_p_values = 10 # Number of different p values to try
max_p = 6000 # Maximum value of p to try
p_values = np.linspace(100, max_p, num_p_values, dtype=int)

training_errors = []
validation_errors = []
```

```
•[162]: # Splitting Training data
X_train_split, X_validation, labels_train_split, labels_val = train_test_split(X_train, labels_train, test_size=0.2, random_state=42)
```

```
[161]: for p in p_values:
    # Generate random matrix G with Gaussian entries
    G = np.random.normal(0, np.sqrt(0.005), size=(p, X_train_split.shape[1]))

    # Generate random vector b with uniform entries
    b = np.random.uniform(0, 2*np.pi, size=p)

    # Transform the training and validation data
    X_train_transformed = np.cos(G @ X_train_split.T + b[:, np.newaxis])
    X_val_transformed = np.cos(G @ X_validation.T + b[:, np.newaxis])

    # One-hot encode the training labels
    Y_train_encode = np.eye(num_classes)[labels_train_split]

    # Train the classifier
    W_hat = train(X_train_transformed.T, Y_train_encode, lambda_)

    # Predict labels for training and validation data
    predicted_labels_train = predict(W_hat, X_train_transformed.T)
    predicted_labels_val = predict(W_hat, X_val_transformed.T)

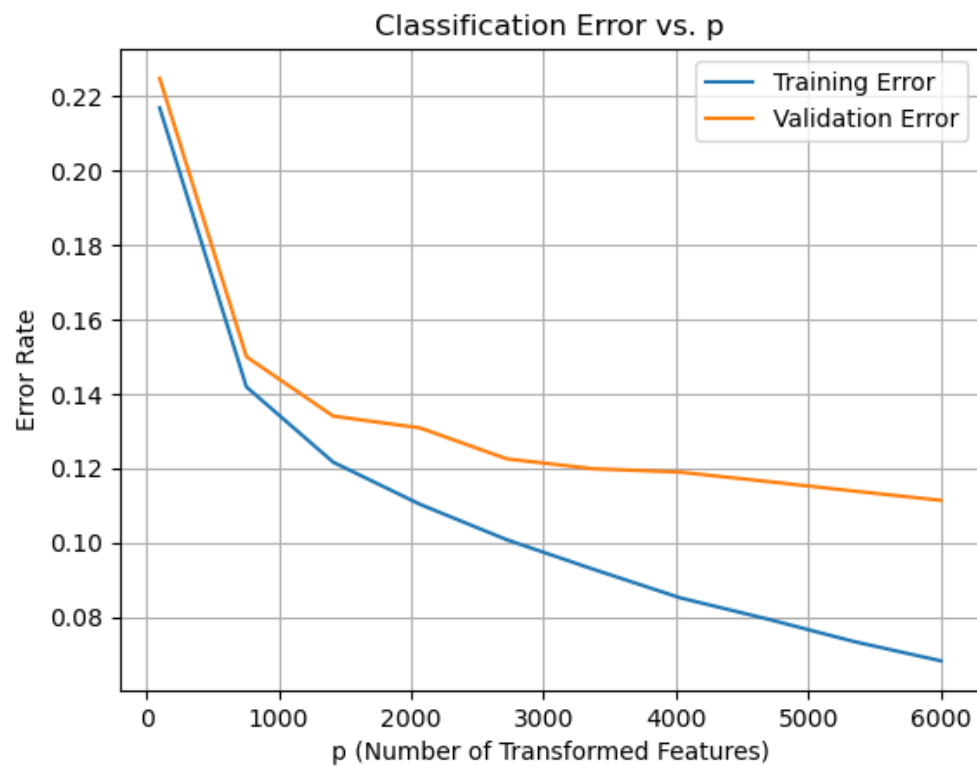
    # Calculate training and validation errors
    training_error = evaluate_prediction(predicted_labels_train, labels_train_split)
    validation_error = evaluate_prediction(predicted_labels_val, labels_val)

    training_errors.append(training_error)
    validation_errors.append(validation_error)
```

Plotting  $\hat{W}^p$  vs  $p$

```
[151]: plt.grid(True)
plt.plot(p_values, training_errors, label='Training Error')
plt.plot(p_values, validation_errors, label='Validation Error')
plt.xlabel('p (Number of Transformed Features)')
plt.ylabel('Error Rate')
plt.title('Classification Error vs. p')
plt.legend()
```

```
[151]: <matplotlib.legend.Legend at 0x199977f55d0>
```



## Q2d. Computing the confidence interval

Transform the test data using the optimal  $p$

```
[125]: p_optimal = p_values[np.argmin(validation_errors)]
print("p_optimal: ", p_optimal)

p_optimal: 6000

•[126]: G_optimal = np.random.normal(0, np.sqrt(0.005), size=(p_optimal, X_train.shape[1]))
b_optimal = np.random.uniform(0, 2*np.pi, size=p_optimal)
X_test_transformed = np.cos(G_optimal @ X_test.T + b_optimal[:, np.newaxis])

•[127]: # Train the classifier using  $\hat{p}$  and Predict for test data
W_hat_optimal = train(X_train_transformed.T, Y_train_encode, lambda_)
predicted_labels_test_optimal = predict(W_hat_optimal, X_test_transformed.T)
print("predicted_labels_test_optimal: ", predicted_labels_test_optimal)

# Calculate the classification test error  $\epsilon_{\text{test}}(\hat{f})$ 
test_error_optimal = evaluate_prediction(predicted_labels_test_optimal, labels_test)

predicted_labels_test_optimal: [9 2 9 ... 9 6 2]

•[128]: confidence_level = 0.95
delta = 1 - confidence_level
a = 0 # Min of true classification error (0%)
b = 1 # Max of true classification error (100%)
m = len(labels_test) # Number of test examples
```

Compute the confidence interval using Hoeffding's inequality

```
[153]: confidence_interval = np.sqrt(((b - a)**2 * np.log(2/delta)) / (2 * m))
print("confidence_interval: ", confidence_interval)

confidence_interval: [0.01355542 0.00932347 0.0706512 ... 0.03441607 0.01300412 0.03359333]

[154]: lower_bound = test_error_optimal - confidence_interval
upper_bound = test_error_optimal + confidence_interval

print("lower_bound: ", lower_bound)
print("upper_bound: ", upper_bound)
print("Classification Test Error: ", test_error_optimal)
print("confidence_level : ", confidence_level)

lower_bound: [0.88194458 0.88617653 0.8248488 ... 0.86108393 0.88249588 0.86190667]
upper_bound: [0.90905542 0.90482347 0.9661512 ... 0.92991607 0.90850412 0.92909333]
Classification Test Error: 0.8955
confidence_level : 0.95
```